
Replication and Predictable-Drift Diagnostics for FTS-Diffusion

Anonymous Author(s)

Affiliation

Address

email

Abstract

We revisit FTS-Diffusion (ICLR 2024), a financial time-series generator built from scale-invariant subsequence clustering (SISC), pattern-conditioned diffusion, and learned pattern evolution, and try to reproduce it from the released code. In our runs the released sampling pipeline is near-degenerate: synthetic S&P 500 paths are one repeated motif tiled into an almost perfectly linear price ramp (linear-fit $R^2 \approx 1.0$), independent runs are nearly identical, and synthetic return volatility is about $20\times$ too small. We trace this to three choices in the released sampling code: a deterministic argmax rollout of the pattern chain, a one-unit decoder hidden state, and a commented-out magnitude normalization. We are explicit about confounds we cannot rule out: no checkpoints are shipped, and our training budget is reduced. Consistent with degeneracy, a continuous-rollout Train-on-Augmentation-Test-on-Real (TATR) setup appears to improve S&P 500 forecasting by up to 85%, yet the same protocol never beats its baseline on GOOG or ZC=F. This non-transfer is what a coincidental scale match would produce, not useful augmentation. On the released artifacts the authors' own Kolmogorov–Smirnov and Anderson–Darling return tests cleanly separate synthetic from real in essentially every window, so a simple return check detects the degeneracy that the downstream MAPE misses. We also do not reproduce the Appendix B.3 SISC toy targets. As a predeclared extension we add a predictable-drift diagnostic that calibrates on real assets, tracks exploitable artificial alpha in a 500-seed experiment, and improves a 1000-seed augmentation-selection benchmark. We frame these as observations from one reproduction attempt, not a verdict on the method.

1 Introduction

Deep generative models are attractive for finance because historical market data are scarce, noisy, nonstationary, and not experimentally reproducible. Financial returns also exhibit heavy tails, weak raw-return autocorrelation, volatility clustering, and recurring shapes at irregular time scales [Cont, 2001]. Sequence generators try to preserve temporal structure with adversarial, optimal-transport, and score-based objectives [Yoon et al., 2019, Xu et al., 2020, Tashiro et al., 2021]; financial generators additionally need to respect stylized facts and economically meaningful dependence [Wiese et al., 2020]. FTS-Diffusion [Huang et al., 2024] addresses this by discovering variable-length motifs and generating new trajectories motif by motif.

This is a replication and diagnostic-extension study, not a new generator. We test whether the published downstream gains are robust to rollout and evaluation details, whether the SISC toy-data targets are reproducible, and whether a simple predictable-drift statistic catches artificial alpha that marginal stylized facts miss. The architecture is well motivated, but the released sampling pipeline, in our hands, produces degenerate artifacts: the synthetic series collapses to a near-linear price ramp built from one repeated motif (Section 4). Under that lens the apparent S&P 500 gain looks less like

robust augmentation than a coincidence that does not transfer across assets. We cannot fully separate this from a known confound: no trained checkpoints are released, and our training budget is shorter than the released defaults.

We therefore extend the evaluation with a predictable-drift diagnostic. Under a fixed information vector

$$\omega_t = (1, r_t, r_{t-1}, |r_t|, |r_{t-1}|)^\top, \quad (1)$$

we measure whether the next return r_{t+1} has a nonzero linear conditional mean under a bounded test class. This is not a proof of market efficiency or absence of arbitrage. It is a finite, reproducible check for an important failure mode: a synthetic generator can match tails and volatility clustering while introducing a simple, exploitable conditional drift.

Contributions. We make five contributions. (i) A code-level diagnosis of the released sampling pipeline (Section 4): three specific choices that drive the synthetic series to a near-degenerate linear ramp, plus the underspecifications we had to resolve. (ii) A six-seed S&P 500 TATR study, with the cross-asset observation that the continuous-rollout gain does not transfer to GOOG or corn futures (ZC=F), which we read as a scale coincidence. (iii) A return-distribution check showing that the authors’ own Kolmogorov–Smirnov and Anderson–Darling tests detect the degenerate output that the downstream MAPE still rewards. (iv) A documented Appendix B.3 SISC toy-data gap. (v) A fixed predictable-drift diagnostic, calibrated on real assets, applied to ordered FTS-Diffusion samples, validated by a 500-seed predictive-alpha experiment, and used in a 1000-seed augmentation-selection benchmark. Throughout, we separate claims about the released artifacts from what would need the authors’ checkpoints to settle.

2 Background: FTS-Diffusion

FTS-Diffusion decomposes financial time-series generation into three modules [Huang et al., 2024]. Scale-invariant subsequence clustering (SISC) segments a univariate series into variable-length motifs using dynamic time warping (DTW) [Sakoe and Chiba, 1978] and DTW barycenter-style centroid updates [Petitjean et al., 2011]. Each segment is summarized by a pattern identity p_m , a duration scale α_m , and a magnitude scale β_m . A pattern-conditioned diffusion model then generates concrete segments from motif conditions using a DDPM-style denoising objective [Ho et al., 2020].

The long-horizon behavior is controlled by the pattern-evolution module, which learns

$$(p_m, \alpha_m, \beta_m) \mapsto (p_{m+1}, \alpha_{m+1}, \beta_{m+1}), \quad (2)$$

with a classification loss for the next pattern and regression losses for the next duration and magnitude scales. This separation is useful, but it also creates a sensitive interface: downstream results depend on how the learned state chain is initialized, reset, and sliced into synthetic training blocks.

3 Replication protocol

Assets and splits. The main datasets are daily Yahoo Finance close-price histories for S&P 500, GOOG, and ZC=F. The S&P 500 setup spans 1980–2020 and yields 10,088 trading days and 2,282 post-SISC segments with $K = 14$. GOOG spans 2004–2020 with 3,776 trading days and 733 post-SISC segments with $K = 11$. ZC=F spans 2000–2020 with 4,764 trading days and 629 post-SISC segments with $K = 11$. We follow a chronological train/test design: motif extraction, fitted generator components, standardizers, and downstream predictors are fit only on their allowed training windows, and final scores are computed on held-out real data.

Downstream benchmark. The main downstream benchmark is TATR: Train on Augmentation, Test on Real. The held-out real slice is divided into an initialization portion and a final real test portion. A one-day-ahead LSTM forecaster is trained on the real initialization portion plus Y synthetic trading years, where one synthetic year has 252 days and Y is swept across a fixed grid. The score is mean absolute percentage error (MAPE) on the final real test portion. We use paper-like LSTM settings: one-day ahead, window size 64, hidden size 32, two layers, MAE training loss, Adam with learning rate 10^{-2} , and 100 epochs for the primary S&P 500 runs.

84 **Synthetic rollout variants.** We compare three ways to produce the synthetic augmentation path.

- 85 • **Independent fixed blocks:** generate independent 252-day blocks from the same fixed initial
86 state.
- 87 • **Continuous chunked:** generate one continuous trajectory with the pattern-evolution chain,
88 then split it into 252-day chunks for downstream training.
- 89 • **Continuous with scaler refit:** generate a continuous trajectory and refit the downstream
90 scaler on the augmented training series.

91 These protocols use the same broad data budget but different Markov-chain connectivity. The
92 distinction matters because the pattern-evolution module is the only component that carries long-
93 horizon state across generated motifs.

94 **Compute.** The primary six-seed S&P 500 downstream batch used stored SISC artifacts and PG-
95 M/PEM checkpoints and recorded 10,478 seconds on an Apple M4 laptop with 10 CPU cores and
96 16 GB memory. The lightweight drift diagnostics used the same machine: the ordered FTS audit
97 recorded 629.82 seconds with a four-thread cap, the 500-seed predictive-alpha script recorded 2.83
98 seconds, and the 1000-seed selection benchmark recorded 32.69 seconds. These times exclude earlier
99 exploratory runs not used for the reported tables.

100 4 Degeneracy of the released generation pipeline

101 We reproduce FTS-Diffusion end to end from the released code. Before downstream scores, we report
102 what the generator actually produces, since this conditions every later interpretation. In our runs the
103 released sampling path is near-degenerate: the synthetic price series is a single repeated motif tiled
104 into an almost perfectly straight line. We give three jointly responsible choices, with file references
105 into our working copy (fts-diffusion-ref/), then show the degeneracy on the saved trajectories
106 and the return distribution. These are observations about the released artifacts under our setup, not
107 proofs about the method; Section 4.5 lists the confounds we cannot rule out.

108 4.1 Three choices in the released sampling code

109 **(1) Deterministic argmax rollout of the pattern chain.** The pattern-evolution module (PEM) is
110 rolled out by taking `argmax` for the next pattern and the next length, and the raw regression mean for
111 the next magnitude, with no sampling (fts-diffusion-ref/models/sampling.py, lines 30, 33,
112 34):

```
113 pred_pattern = torch.argmax(probabilities_pattern, dim=1).unsqueeze(0)  
114 pred_length = torch.argmax(probabilities_length, dim=1).unsqueeze(0) + l_min  
115 pred_magnitude = pred[:, n_patterns+range_length:].float()
```

118 An autoregressive map $\text{state}_{m+1} = f(\text{state}_m)$ that is deterministic must, iterated from a fixed initial
119 state, converge to a fixed point or a short cycle. The rollout policy (greedy vs. sampling) is not
120 specified in the paper, so this is a defensible reading of the release, but it removes the stochasticity
121 the Markov-chain description suggests.

122 **(2) One-unit decoder hidden state.** The Scaling-AE decoder is an LSTM whose hidden di-
123 mension is fed from `pgm_params['sae_hidden_dim']`, set to 1 in the released configuration
124 (`model_params.py`; the decoder LSTM is defined in `scaling_autoencoder.py`). With a single
125 hidden unit the decoder has almost no capacity, and we observe it collapsing to a fixed output that
126 ignores both the conditioning pattern and the diffusion latent.

127 **(3) Commented-out magnitude normalization.** The sampling step that would rescale the gener-
128 ated shape to unit range before multiplying by the magnitude β is commented out (`sampling.py`,
129 lines 53–55), so the emitted segment has magnitude $\beta \times$ (whatever range the decoder produced)
130 rather than β . In our runs that range is on the order of 0.07, far below the intended scale, which
131 removes the magnitude information β was meant to carry.

4.2 What the saved trajectories look like

Figure 1 contrasts a real S&P 500 path with one released-pipeline synthetic path. The synthetic series is visually a straight line; a least-squares linear fit gives $R^2 \approx 0.99999$. Three quantitative observations accompany the figure. First, independent runs of a configuration are nearly identical, consistent with a deterministic rollout. Second, the PEM state sequence collapses (Table 1): on S&P 500 a single pattern is used in 99.6% of segments (a fixed point), while GOOG and ZC=F settle into period-2 cycles. Third, inspecting the trained PGM decoder pattern by pattern, the decoder output is near-constant across all patterns (cross-pattern standard deviation < 0.001): an up-trend pattern and a down-trend pattern yield essentially the same segment. The net effect is that one tiny upward-creeping segment is tiled across the horizon, which is exactly the linear ramp in Figure 1.




Figure 1: Released-pipeline output on S&P 500. (a) Real close prices show rich structure. (b) One synthetic continuous-rollout path is almost perfectly linear (linear-fit $R^2 \approx 0.99999$). (c) Daily-return density (log scale): synthetic returns concentrate near a single value with standard deviation about $20\times$ smaller than real returns. We show one representative run; runs are nearly identical.

Table 1: PEM state-sequence collapse under the deterministic rollout, from saved `run*_states.npy`. “Dominant pattern share” is the fraction of generated segments using the most frequent pattern.

Asset (K)	Segments	Dominant pattern share	Regime
S&P 500 (14)	2520	99.6%	fixed point
GOOG (11)	2520	50.0%	period-2 cycle
ZC=F (11)	1867	50.0%	period-2 cycle

4.3 A return-distribution check detects the degeneracy

A natural and direct check is the authors’ own return-distribution test. Following their `distribution_tests` routine, we form simple returns and compare real against synthetic with a Kolmogorov–Smirnov (KS) two-sample test and an Anderson–Darling (AD) k -sample test, where a higher p -value means the two return samples are harder to distinguish. The paper reports values consistent with indistinguishability (KS around 0.32, AD around 0.12). On the released artifacts the test instead separates the two cleanly (Table 2): across 300 windows per asset and protocol the synthetic and real returns differ in essentially every window, with mean KS p -values at the 10^{-4} to 10^{-7} level and no window passing at $p > 0.05$. The mechanism is visible in the return moments: the released synthetic series has 4 to $60\times$ too little return volatility, because a near-linear ramp has almost no daily variation. The point is not that we failed to match a number: a degenerate generator is expected to fail a distributional test. It is that this simple return check is *not* blind to the degeneracy, in contrast to the downstream price-MAPE of Section 5 that the same degenerate output can still improve.

4.4 Underspecifications we had to resolve

Several choices that matter for the result are not pinned down by the paper. Table 3 lists the ones we encountered, separating the three that are load-bearing for the degeneracy above from the compute-driven confounds and the remaining setup choices. We kept the released values wherever they existed, so that the run reflects the artifact as published rather than our preferences. We note this is an audit aid, not an accusation: underspecification is common, but here it coincides with the failure modes.

Table 2: Return-distribution check on the released artifacts. “Mean KS p ” and “% indist.” are over 300 windows of length 60; higher is more similar (the paper reports KS around 0.32). “Real/Syn. sd” are daily-return standard deviations. `single` is a continuous rollout, `authors` the independent-block protocol.

Asset	Protocol	Real sd	Syn. sd	Mean KS p	% indist.
S&P 500	authors	0.0109	0.0045	9×10^{-4}	0
S&P 500	single	0.0109	0.00046	4×10^{-7}	0
GOOG	authors	0.0184	0.0026	6×10^{-5}	0
GOOG	single	0.0184	0.00029	4×10^{-7}	0
ZC=F	authors	0.0181	0.0028	2×10^{-5}	0
ZC=F	single	0.0181	0.00041	1×10^{-6}	0

Table 3: Paper-vs-release underspecifications we resolved. “Released” marks a value shipped in the released configuration that we kept; “ours” marks a choice we had to make.

Design choice	In paper?	Value used	Note
Scaling-AE decoder hidden dim	no	1 (released)	load-bearing
PEM rollout policy	no	argmax (released)	load-bearing
Magnitude normalization at sampling	implied	commented out (released)	load-bearing
PGM / PEM training epochs	no	30 / 60 (ours; released 400 / 1000)	confound
PCDM diffusion steps	no	30 (ours; released 100)	confound
SISC k -means seed	no	42 (released)	shared seed
GOOG / ZC=F downstream split	S&P only	62.5% train (ours)	cross-asset parity
Initial segment for generation	no	first test segment / random (ours)	protocol switch

162 4.5 Confounds and what we cannot rule out

163 We are explicit that the evidence above concerns the released artifacts as configured and as trained
164 in our hands, not the method in principle. Two confounds prevent a stronger claim. First, no
165 trained checkpoints are shipped with the release, so every reproduction retrains from the released
166 scripts and we cannot compare against the authors’ own generations. Second, our training budget
167 is reduced relative to the released defaults (PGM 30 vs. 400 epochs, PEM 60 vs. 1000, PCDM 30
168 vs. 100 steps), a compute-driven choice that a reviewer can reasonably challenge. The degeneracy is
169 therefore consistent with at least three non-exclusive explanations that we have not fully separated:
170 undertraining, an unlucky single seed, and the one-unit decoder capacity. A full-budget retrain across
171 seeds would falsify or confirm the undertraining hypothesis; we flag it as the natural next experiment
172 and do not assert a method-level conclusion without it. Given how unusual a one-unit decoder hidden
173 state is for a published generative model, it would also be reasonable to ask the authors directly
174 whether that value, and the commented-out magnitude normalization, correspond to the configuration
175 behind the paper’s headline results.

176 5 Replication results

177 5.1 TATR depends on the synthetic-series protocol

178 The original paper reports that appending 100 synthetic years reduces one-day-ahead MAPE by 17.9%,
179 15.3%, and 17.4% on S&P 500, GOOG, and ZC=F, respectively. Our current S&P 500 evidence has
180 three layers. First, the stored result matrices do not support that trend for the documented authors
181 independent-block protocol: the stored S&P 500 authors curve finishes 208.4% worse than its
182 baseline, while the `single` continuous-trajectory curve finishes 71.6% better. Second, a one-seed
183 protocol search found that continuous-rollout variants finish 76.5–85.4% better, whereas independent
184 fixed blocks finish 96.5% worse. Third, Table 4 reports the current primary six-seed long batch.

185 In that long batch, continuous protocols reproduce strong paper-like reductions in downstream MAPE,
186 but independent fixed blocks do not. The conclusion is not driven by a single lucky seed: the final
187 continuous-chunked changes range from −86.8% to −55.1%, while the final independent-block

changes range from +37.1% to +296.0%. The contrast is also visible in Figure 2: continuous trajectories eventually cross below the real-only baseline, while independent blocks worsen as augmentation increases.

Table 4: Six-seed S&P 500 TATR summary. Percentages are relative to each protocol’s no-augmentation baseline; lower MAPE is better. Ranges are seed-wise final changes.

Protocol	Baseline	Best change	Final change	Final range
Continuous chunked, 1-day	0.0589	−84.9%	−65.2%	[−86.8, −55.1]%
Continuous + scaler refit, 1-day	0.0630	−85.4%	−68.2%	[−84.3, −30.1]%
Independent fixed blocks, 1-day	0.0646	0.0%	+193.2%	[+37.1, +296.0]%
Continuous chunked, 5-day	0.0810	−85.2%	−77.5%	[−85.9, −58.5]%




Figure 2: TATR outcome changes sharply with the synthetic-series protocol. Continuous trajectories can look paper-like, while independent blocks worsen with augmentation. The dashed line is the original S&P 500 100-year TATR target. The shaded bands show min/max across seeds.

The mechanism is visible in the synthetic price levels. Continuous trajectories start near 1253 and end near 7085 (mean level 4138); independent blocks start from the same level but repeatedly reset, ending near 1249 (mean 1214). The continuous protocol thus drifts its near-linear ramp through the scale of the upward-trending S&P 500 test period, so percentage error can fall simply because synthetic and real price levels overlap, not because the dynamics are informative. Independent blocks stay near the initial level, the overlap never happens, and error grows. Read with Section 4, this is a scale effect, not useful augmentation. The cross-asset evidence below is the cleaner test.

TMTR controls make the replication conclusion more conservative. In the six-seed long batch, the reference TMTR curve reaches a best mean improvement of 26.8% but finishes 359.7% worse; a continuous offset-30 control reaches a best improvement of 50.0% and finishes 17.8% better; a continuous offset-50 control finishes 752.0% worse. Thus the evidence is not “synthetic data always helps.” It is narrower: a particular continuous TATR price-level rollout can create large apparent gains, while other augmentation and evaluation choices do not support the same conclusion.


We make no claim of misconduct; the available generation details do not pin down the downstream result uniquely. We note only that, once the generator is seen to be degenerate, the most economical reading of the S&P 500 gain is a scale coincidence rather than augmentation that works under the right protocol. Either way, a reproducible benchmark should treat the Markov-chain rollout, initialization, block slicing, scaler fitting, and warm-up windows as first-class parameters, and should not rest a positive conclusion on downstream MAPE for a single trending asset.

GOOG and ZC=F test the coincidence reading most directly. We regenerated the stated real-asset SISC and architecture artifacts for both with $K = 11$, $l_{\min} = 10$, and $l_{\max} = 21$. A strict author-faithful downstream TATR rerun is blocked by the released S&P-style split logic, which consumes $252 \times 5 = 1260$ held-out initialization points whereas the GOOG and ZC=F 80/20 held-out windows contain 744 and 963 points, so we adapted the split while keeping the proportions comparable. This changes the absolute MAPE level but not the sign of the augmentation effect, since every percentage is computed within the same split relative to that protocol’s own no-augmentation baseline. Under that within-split comparison the effect is negative on both assets and for both protocol families. The independent-block (authors) protocol finishes 1003.1% worse on GOOG and 75.1% worse on ZC=F. Crucially, the *continuous* rollout that appears to help on S&P 500 does not transfer: its best point over the whole sweep equals the no-augmentation baseline (it never beats it), and it finishes about 106% worse on GOOG and 179% to 431% worse on ZC=F depending on the variant. A protocol that helped because the generator carried useful structure should not flip sign across assets; the continuous gain confined to the one strongly trending asset is the signature of a scale coincidence, not of augmentation value. We therefore treat GOOG and ZC=F as covered for SISC/architecture setup and real-asset diagnostics rather than author-faithful downstream replications, but the within-split sign does not depend on author-faithfulness.

5.2 SISC toy-data replication gap

We also stress-tested SISC on the Appendix B.3 simulated-data setting. The paper reports one-pattern per-unit DTW/Jaccard of 0.009/0.938 and multi-pattern average per-unit DTW/Jaccard of 0.01/0.784. On our derived one-pattern runs, the best available DTW is 0.0202 and the best boundary-Jaccard-at-tolerance-2 proxy is 0.8693. Across 18 successful multi-pattern runs varying seed, $l_{\max}$, and maximum iterations, our best average per-unit DTW was 0.0321, and our best author-style interval IoU was 0.6418. Figure 3 summarizes the multi-pattern gap.

Appendix B.3 Replication Gap



B.3 targets come from the paper. Local values use the shipped toy data and reference SISC with runtime fixes.

Figure 3: Appendix B.3 SISC replication gap on the toy data. The best replication runs do not reach the reported centroid or segmentation targets.

The gap persisted when increasing SISC iterations from 10 to 20 and when changing $l_{\max}$ from 20 to 21. The best strict boundary Jaccard at tolerance 2 was 0.3608, and the best reference-code-style interval IoU was 0.6418. The best centroid metric remained about $3.2\times$ the reported target. What failed here is reproducibility of the exact toy benchmark, not the general idea of SISC: the released tree does not expose the exact one-pattern construction, synthetic generator, standard pattern arrays, seed state, or metric implementation used to produce Appendix B.3. The broader replication lesson is that helper defaults matter: warm-up windows, augmentation scales, plotting conventions, and interval-overlap metrics can all change the apparent result if they are not audited.

6 Predictable-drift extension

6.1 Diagnostic

Stylized facts are necessary but not sufficient for finance. A generator can reproduce heavy tails and volatility clustering while encoding a simple conditional mean, such as $r_{t+1} = \rho r_t + \epsilon_{t+1}$. Marginal tests may not detect this, but a trading signal using r_t can. We therefore treat the drift check as a fixed, predeclared diagnostic rather than a learned discriminator: the chronological split rule, feature map, train-only standardization, null construction, and companion stylized-fact summaries are fixed before evaluating synthetic outcomes.

Motivated by the martingale-difference intuition behind weak-form market efficiency [Fama, 1970], we fix a small information set

$$\omega_t = (1, r_t, r_{t-1}, |r_t|, |r_{t-1}|)^\top \quad (3)$$

and a bounded linear class

$$g_b(\mathcal{F}_t) = b^\top \omega_t, \quad \|b\|_2 \leq 1. \quad (4)$$

The population predictable-drift violation is

$$\delta = \sup_{\|b\|_2 \leq 1} |\mathbb{E}[r_{t+1} b^\top \omega_t]|. \quad (5)$$

Let $m = \mathbb{E}[r_{t+1} \omega_t]$. By Cauchy-Schwarz, Equation 5 has the closed form

$$\delta = \|m\|_2. \quad (6)$$

The empirical estimator is therefore

$$\hat{\delta} = \left\| \frac{1}{T-2} \sum_{t=2}^{T-1} r_{t+1} \omega_t \right\|_2. \quad (7)$$

The statistic decomposes into interpretable components: unconditional drift, first-lag return predictability, second-lag return predictability, and two volatility-conditioned mean terms. In our implementation, returns are standardized using training-split mean and standard deviation only; the constant feature remains equal to one. We compare real data, synthetic data, and a block-shuffled null that preserves the return multiset and some local dependence while disrupting longer chronological alignment.

6.2 Real-asset calibration

Before using $\hat{\delta}$ as a rejection statistic, we calibrate it on real held-out assets. For S&P 500, GOOG, and ZC=F, we start from raw close-price CSVs, use the first 80% of simple returns for standardization, and evaluate the last 20%. The null distribution uses 500 replications of 21-day block shuffling on the held-out returns. Table 5 shows that all three real held-out paths lie inside their asset-specific 5–95% null bands, while retaining heavy tails and volatility clustering. This is a useful sanity check: the diagnostic is not simply rejecting real financial data because it is non-Gaussian.

Table 5: Real-asset predictable-drift calibration. The null band is the 5–95% range from 500 block-shuffled held-out paths; the final columns report stylized-fact context.

Asset	Test obs.	Real $\hat{\delta}$	Null 5–95%	ACF($ r $)	Ex. kurt.
S&P 500	2018	0.0297	[0.0222, 0.0303]	0.2165	3.2110
GOOG	755	0.0575	[0.0403, 0.0629]	0.0935	6.1142
ZC=F	953	0.0173	[0.0150, 0.0259]	0.1250	2.0247

6.3 Controlled predictive-alpha test

We next test whether larger $\hat{\delta}$ corresponds to exploitable predictability. In each of 500 seeds, we simulate a zero-mean heavy-tailed GARCH-like path, keep the evaluation-period magnitude path

Table 6: Controlled predictive-alpha test over 500 seeds. Magnitude paths are fixed within seed; only sign dependence changes.

Sample	p	$\hat{\delta}$	OOS corr.	Dir. acc.	Sharpe (sd)
Real evaluation path	—	0.0976	0.0004	0.5025	0.07 (0.77)
IID signs	—	0.0941	0.0036	0.5020	0.09 (0.74)
Markov signs	0.65	0.2088	0.1482	0.5968	2.33 (0.96)
Markov signs	0.75	0.3435	0.2786	0.6952	4.80 (1.05)
Markov signs	0.85	0.5126	0.4112	0.7976	7.72 (1.18)
Markov signs	0.90	0.6164	0.4811	0.8483	9.43 (1.33)




Figure 4: Drift versus no-drift evidence from the larger controlled experiments. The diagnostic tracks held-out artificial alpha, separates hidden-drift sign processes from real and iid controls, and improves augmentation selection by avoiding null violations and high-persistence candidates.

272 $|r_t|$ fixed, and replace only the signs. The iid-sign control has no intentional sign predictability;
 273 Markov-sign samples have sign persistence p . We fit a ridge linear predictor on the first half of each
 274 evaluation path using the same feature map ω_t , then evaluate one-step prediction and a sign strategy
 275 on the second half.

276 Table 6 and Figure 4 show the result. The iid-sign and real controls have near-zero out-of-sample
 277 correlation, near-50% directional accuracy, and near-zero strategy Sharpe. At $p = 0.85$, the hidden-
 278 drift sample has mean $\hat{\delta} = 0.5126$, out-of-sample correlation 0.4112, directional accuracy 79.76%,
 279 and mean annualized Sharpe 7.72 with seed standard deviation 1.18. Across all controlled samples,
 280 pooled $\hat{\delta}$ and out-of-sample Sharpe have correlation 0.6921. Thus the diagnostic is not merely a
 281 reporting statistic; it is aligned with artificial alpha.

282 6.4 Ordered FTS audit and incorporated selection

283 The FTS-Diffusion use case is direct: report Equation 7 on ordered samples from each rollout protocol,
 284 and do not select a protocol only because it improves TATR if its $\hat{\delta}$ lies outside the real/null band.

Table 7: Ordered S&P 500 FTS-Diffusion drift audit over six seeds and 25,200 generated prices per seed. The cutoff is the S&P 500 held-out block-shuffle null 95th percentile, $q_{.95} = 0.0303$.

Protocol	$\hat{\delta}$	$\hat{\delta}/q_{.95}$	Violations	Price std.
Continuous rollout	0.0263	0.87	0/6	1701.5
Independent fixed blocks	0.0370	1.22	6/6	24.5

Table 8: Drift-incorporated augmentation selection over 1000 seeds. Lower normalized MSE is better. High- p is the fraction of selected Markov-sign candidates with persistence $p \geq 0.75$. The oracle uses held-out test MSE and is only a reference upper bound.

Selector	Null viol.	High- p	Styl. score	Test NMSE	Excess
Oracle test-MSE	41.5%	0.5%	0.2135	0.9979	-0.0319
Drift-penalized	34.6%	0.2%	0.0520	1.0087	-0.0211
Stylized-only	76.0%	39.3%	0.0344	1.0452	+0.0154
Random	84.8%	50.2%	0.3106	1.0579	+0.0281

We regenerated six S&P 500 synthetic paths per protocol from stored FTS-Diffusion checkpoints, using 100 ordered 252-day blocks per seed and excluding artificial transitions between independent blocks. Table 7 shows that continuous rollouts stay inside the S&P 500 null band in all six seeds, while independent fixed blocks violate it in all six. The audit therefore does not reject the continuous rollout that produces paper-like TATR gains; it identifies the independent reset protocol as both downstream-harmful and predictably drifted under this statistic.

We next test whether the same principle helps choose augmentations in a controlled benchmark. For each of 1000 seeds, all candidate augmentations shared the same magnitude path but had different sign dynamics. A marginal stylized selector minimized a distance over volatility, tail, quantile, and absolute-return autocorrelation summaries; the drift-incorporated selector minimized

$$d_{\text{styl}} + 0.25 \left[\hat{\delta}/q_{.95}^{\text{real}} - 1 \right]_+^2, \quad (8)$$

where $q_{.95}^{\text{real}}$ is the real block-shuffle null cutoff. The selected augmentation trained the same one-step ridge predictor on a 252-day real initialization plus synthetic returns, and performance was measured on a final held-out real path that the drift-penalized selector does not see.

Table 8 shows that incorporating the diagnostic changes the selected augmentation, not just the reporting. Relative to stylized-only selection, the drift-penalized selector reduces high-persistence choices from 39.3% to 0.2%, cuts real-null violations from 76.0% to 34.6%, and lowers held-out normalized MSE by 0.0365 on average, with a 60.8% paired win rate. If the statistic is incorporated during training rather than selection, use it as a validation-selected penalty

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{FTS}} + \lambda \hat{\delta}_{\text{synthetic}}^2. \quad (9)$$

The penalty must be computed on ordered return sequences. Applying it across randomly shuffled mini-batches would destroy the time ordering and should only be logged as a proxy, not as the final diagnostic.

7 Limitations and broader impact

Our replication is constrained by the information and compute available for this study. Most importantly, the degeneracy analysis of Section 4 concerns the released artifacts as configured and as trained in our hands: no checkpoints are shipped, and our training budget is reduced relative to the released defaults. We have not run the full-budget, multi-seed retrain that would separate undertraining from an architectural cause, so we deliberately phrase those findings as observations about the release rather than conclusions about the method. Some runs use reference checkpoints rather than a full fresh retraining of all modules, and the ordered FTS-Diffusion drift audit is limited to the S&P 500 checkpoints and the fixed diagnostic in Equation 7. Train-time validation selection of λ remains to be run before claiming an optimized drift-regularized generator. GOOG and ZC=F have

316 shorter histories than S&P 500, so the downstream split used for S&P 500 is not directly portable
317 without changing the question being asked; we mitigate this by comparing within split, relative to
318 each protocol’s own baseline.

319 The predictable-drift diagnostic is deliberately narrow. It tests a fixed finite-dimensional feature
320 map, not every admissible trading strategy, and it is not a proof of no-arbitrage. Its strength is
321 reproducibility: the statistic is simple, closed-form, and hard to tune after seeing the result.

322 Synthetic financial data can be beneficial for stress testing, augmentation, and reproducibility, but it
323 can also create misleading evidence if artificial predictability is mistaken for real market structure.
324 Any release of synthetic paths should clearly label them as generated data, document the training
325 period and generation protocol, and avoid implying that downstream forecasting gains translate into
326 tradable performance.

327 8 Conclusion

328 FTS-Diffusion has an appealing inductive bias for financial time series: variable-length motifs,
329 pattern-conditioned diffusion, and learned motif evolution. Our reproduction attempt, however, did
330 not yield a working generator from the released artifacts. In our hands the released sampling pipeline
331 collapses to a near-linear ramp built from one repeated motif, the authors’ own return-distribution
332 test cleanly flags that output as far from real, and the one downstream setting that appears to improve
333 forecasting does so only on a single strongly trending asset and fails to transfer to GOOG or ZC=F.
334 The most economical reading of that pattern is a scale coincidence rather than useful augmentation.
335 We are careful not to overstate this: we cannot yet separate the degeneracy from undertraining without
336 the authors’ checkpoints and full training budget, so we present empirical observations about the
337 release, not a verdict on the method. What the evidence does support is a methodological point.
338 Realistic-looking paths and a favorable downstream MAPE on a trending series are weak evidence on
339 their own. Future synthetic-finance benchmarks should report exact rollout protocols, direct generator
340 diagnostics such as the return-distribution and predictable-drift checks used here, and cross-asset
341 transfer, rather than resting positive conclusions on a single asset.

342 References

- 343 Rama Cont. Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative*
344 *Finance*, 1(2):223–236, 2001.
- 345 Eugene F. Fama. Efficient capital markets: A review of theory and empirical work. *The Journal of*
346 *Finance*, 25(2):383–417, 1970.
- 347 Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In *Advances in*
348 *Neural Information Processing Systems*, volume 33, pages 6840–6851, 2020.
- 349 Hongbin Huang, Minghua Chen, and Xiao Qiao. Generative learning for financial time series with
350 irregular and scale-invariant patterns. In *International Conference on Learning Representations*,
351 2024. URL <https://openreview.net/forum?id=CdjnzWsQax>.
- 352 Francois Petitjean, Alain Ketterlin, and Pierre Gancarski. A global averaging method for dynamic
353 time warping, with applications to clustering. *Pattern Recognition*, 44(3):678–693, 2011.
- 354 Hiroaki Sakoe and Seibi Chiba. Dynamic programming algorithm optimization for spoken word
355 recognition. *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 26(1):43–49, 1978.
- 356 Yusuke Tashiro, Jiaming Song, Yang Song, and Stefano Ermon. CSDI: Conditional score-based
357 diffusion models for probabilistic time series imputation. In *Advances in Neural Information*
358 *Processing Systems*, volume 34, pages 24804–24816, 2021.
- 359 Magnus Wiese, Robert Knobloch, Ralf Korn, and Peter Kretschmer. Quant GANs: Deep generation
360 of financial time series. *Quantitative Finance*, 20(9):1419–1440, 2020.
- 361 Tianlin Xu, Li K. Wenliang, Michael Munn, and Beatrice Acciaio. COT-GAN: Generating sequen-
362 tial data via causal optimal transport. In *Advances in Neural Information Processing Systems*,
363 volume 33, pages 8798–8809, 2020.

364 Jinsung Yoon, Daniel Jarrett, and Mihaela van der Schaar. Time-series generative adversarial networks.
365 In *Advances in Neural Information Processing Systems*, volume 32, 2019.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes].

Justification: The abstract and introduction state that this is a replication and diagnostic-extension study, not a new generator. They distinguish the completed S&P 500 downstream replication, GOOG/ZC=F boundary, real-asset calibration, ordered S&P 500 FTS-Diffusion drift audit, predictive-alpha results, and augmentation-selection benchmark from broader generator-level claims.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes].

Justification: The limitations section discusses artifact dependence, the S&P-only scope of the ordered FTS-Diffusion drift audit, remaining train-time penalty selection, shorter non-S&P histories, and the finite-dimensional nature of the drift diagnostic.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A].

Justification: The paper does not claim a new theorem. The only derivation is the closed-form norm expression for the fixed linear drift statistic, which follows directly from Cauchy-Schwarz and is stated in the method section.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes].

Justification: The replication protocol, key hyperparameters, datasets, stored-output audit, protocol-search-to-long-batch evidence chain, compared rollout protocols, train-only standardization, block-shuffle null calibration, controlled predictive-alpha setup, drift-penalized selection rule, and figure-generation source artifacts are described in the main paper or generated experiment outputs. The central claims are protocol sensitivity and diagnostic value, and both are tied to explicit experimental definitions.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.

- (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No].

Justification: This anonymous PDF submission does not include a public code or data URL. The paper reports the experimental settings, protocols, compute, and generated-result summaries needed to interpret the main claims, but an archival de-anonymized code release is deferred to a public version.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes].

Justification: The main text specifies the data assets, chronological split structure, TATR protocol, LSTM settings, rollout variants, TMTR controls, fixed drift diagnostic choices, ordered FTS audit design, real-asset null construction, 500-seed controlled predictive-alpha setup, and 1000-seed augmentation-selection design.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes].

Justification: The main TATR comparison reports six-seed means and min/max bands, the SISC toy-data claim is based on an 18-run sweep, the real-asset calibration reports block-shuffle 5–95% null bands, the ordered FTS audit reports six-seed null-band violations, the predictive-alpha experiment reports 500-seed means with seed standard deviations for the main Sharpe comparison, and the selection benchmark reports 1000 paired seeds with win rate and normalized-MSE reduction.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes].

Justification: The main paper now reports the hardware class, memory, CPU count, and measured runtimes for the primary six-seed S&P 500 downstream batch, the ordered FTS audit, the 500-seed predictive-alpha diagnostic, and the 1000-seed selection benchmark; generated metadata files record the corresponding elapsed seconds.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes].

Justification: The study uses public financial time-series data and replication artifacts, does not involve human subjects, and discusses risks of misleading synthetic financial data.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes].

Justification: The limitations and broader-impact section discusses beneficial uses such as stress testing and reproducibility, as well as risks from artificial predictability and misleading downstream forecasting claims.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A].

Justification: The submission does not release a high-risk pretrained model or scraped personal dataset. It recommends clear labeling and protocol documentation for synthetic financial paths.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.

- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [No].

Justification: The original FTS-Diffusion paper is cited and the financial data sources are named, but the PDF does not enumerate all upstream code and data licenses or terms of use. This is a limitation of the current anonymous artifact package.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A].

Justification: The submission does not introduce or release a new dataset, pretrained model, benchmark suite, or public software asset. The predictable-drift diagnostic is described as a method and evaluation statistic in the paper.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A].

Justification: The work does not involve crowdsourcing or human-subject experiments.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A].

Justification: The work does not involve human subjects and therefore does not require IRB approval.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A].

Justification: LLMs are not part of the scientific method, model, or diagnostic proposed in this paper.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.