


What Gets Diffused in FTS-Diffusion?


The SISC motif is the condition; the diffusion model learns how real segments vary around that condition.

1. SISC library



Real series is segmented into motifs and each segment receives a pattern label.

2. Condition p




One motif is selected as the conditioning shape for the generator.

3. Add noise




Training adds scheduled Gaussian noise to the residual around that motif.

4. Denoise



Sampling starts from noise and repeatedly predicts the noise to remove.

5. Output segment



The result is decoded, rescaled by magnitude, and stitched into the synthetic path.

Speaker line: the model is not diffusing the pattern library itself. It keeps the pattern visible and denoises a residual/latent segment around it.